

ASSIGNMENT 11.1

1) Transfer data between mysql and hdfs using sqoop import and export

Apache Sqoop is a tool designed to efficiently transfer bulk data between [Hadoop](#) and structured datastores such as relational databases.

```
mysql> create database db1;  
Query OK, 1 row affected (0.05 sec)
```

```
mysql> use db1;  
Database changed  
mysql> create table company  
-> (id int,  
-> name varchar(20),  
-> location varchar(20));  
Query OK, 0 rows affected (0.08 sec)
```

```
mysql> insert into company values(1,'acadgild','banglure');  
Query OK, 1 row affected (0.11 sec)
```

```
mysql> insert into company values(2,'spring','banglure');  
Query OK, 1 row affected (0.00 sec)
```

```
mysql> insert into company values(3,'ibm','cochin');  
Query OK, 1 row affected (0.00 sec)
```

```
mysql> select * from company;
```

```
+-----+-----+-----+  
| id | name | location |  
+-----+-----+-----+  
| 1 | acadgild | banglure |  
| 2 | spring | banglure |  
| 3 | ibm | cochin |  
+-----+-----+-----+
```

```
3 rows in set (0.00 sec)
```

```
mysql> grant all on *.* to 'root@localhost' with grant option;  
Query OK, 0 rows affected (0.00 sec)
```

```
mysql> flush privileges;  
Query OK, 0 rows affected (0.03 sec)
```

```
mysql> commit;  
Query OK, 0 rows affected (0.00 sec)
```

SQOOP IMPORT

```
[acadgild@localhost ~]$ sqoop import --connect jdbc:mysql://localhost/db1
--username 'root' -P --table 'company' --target-dir '/sqoopout1' -m 1
```

Warning: /usr/local/sqoop/./hcatalog does not exist! HCatalog jobs will fail.
Please set \$HCAT_HOME to the root of your HCatalog installation.

Warning: /usr/local/sqoop/./accumulo does not exist! Accumulo imports will fail.
Please set \$ACCUMULO_HOME to the root of your Accumulo installation.

Warning: /usr/local/sqoop/./zookeeper does not exist! Accumulo imports will fail.
Please set \$ZOOKEEPER_HOME to the root of your Zookeeper installation.

2017-11-03 13:33:06,641 INFO [main] sqoop.Sqoop: Running Sqoop version:
1.4.5

Enter password:

2017-11-03 13:33:10,007 INFO [main] manager.MySQLManager: Preparing to use
a MySQL streaming resultset.

2017-11-03 13:33:10,007 INFO [main] tool.CodeGenTool: Beginning code
generation

2017-11-03 13:33:10,641 INFO [main] manager.SqlManager: Executing SQL
statement: SELECT t.* FROM `company` AS t LIMIT 1

2017-11-03 13:33:10,750 INFO [main] manager.SqlManager: Executing SQL
statement: SELECT t.* FROM `company` AS t LIMIT 1

2017-11-03 13:33:10,760 INFO [main] orm.CompilationManager:

HADOOP_MAPRED_HOME is /usr/local/hadoop-2.6.0

Note:

/tmp/sqoop-acadgild/compile/70ff055bd06b630fd4edc566f0d57bc6/company.java
uses or overrides a deprecated API.

Note: Recompile with -Xlint:deprecation for details.

2017-11-03 13:33:14,790 INFO [main] orm.CompilationManager: Writing jar file:
/tmp/sqoop-acadgild/compile/70ff055bd06b630fd4edc566f0d57bc6/company.jar

2017-11-03 13:33:14,876 WARN [main] manager.MySQLManager: It looks like you
are importing from mysql.

2017-11-03 13:33:14,876 WARN [main] manager.MySQLManager: This transfer
can be faster! Use the --direct

2017-11-03 13:33:14,876 WARN [main] manager.MySQLManager: option to
exercise a MySQL-specific fast path.

2017-11-03 13:33:14,876 INFO [main] manager.MySQLManager: Setting zero
DATETIME behavior to convertToNull (mysql)

2017-11-03 13:33:14,886 INFO [main] mapreduce.ImportJobBase: Beginning
import of company

SLF4J: Class path contains multiple SLF4J bindings.

SLF4J: Found binding in

[jar:file:/usr/local/hbase/lib/slf4j-log4j12-1.6.4.jar!/org/slf4j/impl/StaticLoggerBinder
.class]

SLF4J: Found binding in

[jar:file:/usr/local/hadoop-2.6.0/share/hadoop/common/lib/slf4j-log4j12-1.7.5.jar!/o
rg/slf4j/impl/StaticLoggerBinder.class]

SLF4J: See http://www.slf4j.org/codes.html#multiple_bindings for an explanation.

2017-11-03 13:33:15,941 WARN [main] util.NativeCodeLoader: Unable to load
native-hadoop library for your platform... using builtin-java classes where

applicable

2017-11-03 13:33:15,964 INFO [main] Configuration.deprecation: mapred.jar is deprecated. Instead, use mapreduce.job.jar

2017-11-03 13:33:17,083 INFO [main] Configuration.deprecation: mapred.map.tasks is deprecated. Instead, use mapreduce.job.maps

2017-11-03 13:33:17,459 INFO [main] client.RMProxy: Connecting to ResourceManager at /0.0.0.0:8032

2017-11-03 13:33:21,980 INFO [main] db.DBInputFormat: Using read committed transaction isolation

2017-11-03 13:33:22,310 INFO [main] mapreduce.JobSubmitter: number of splits:1

2017-11-03 13:33:22,431 INFO [main] Configuration.deprecation: mapred.job.name is deprecated. Instead, use mapreduce.job.name

2017-11-03 13:33:22,436 INFO [main] Configuration.deprecation: mapred.cache.files.timestamps is deprecated. Instead, use mapreduce.job.cache.files.timestamps

2017-11-03 13:33:22,436 INFO [main] Configuration.deprecation: mapreduce.map.class is deprecated. Instead, use mapreduce.job.map.class

2017-11-03 13:33:22,439 INFO [main] Configuration.deprecation: mapreduce.inputformat.class is deprecated. Instead, use mapreduce.job.inputformat.class

2017-11-03 13:33:22,439 INFO [main] Configuration.deprecation: mapreduce.outputformat.class is deprecated. Instead, use mapreduce.job.outputformat.class

2017-11-03 13:33:22,439 INFO [main] Configuration.deprecation: mapred.output.value.class is deprecated. Instead, use mapreduce.job.output.value.class

2017-11-03 13:33:22,440 INFO [main] Configuration.deprecation: mapred.output.dir is deprecated. Instead, use mapreduce.output.fileoutputformat.outputdir

2017-11-03 13:33:22,440 INFO [main] Configuration.deprecation: mapred.cache.files is deprecated. Instead, use mapreduce.job.cache.files

2017-11-03 13:33:22,440 INFO [main] Configuration.deprecation: mapred.working.dir is deprecated. Instead, use mapreduce.job.working.dir

2017-11-03 13:33:22,455 INFO [main] Configuration.deprecation: mapred.job.classpath.files is deprecated. Instead, use mapreduce.job.classpath.files

2017-11-03 13:33:22,455 INFO [main] Configuration.deprecation: user.name is deprecated. Instead, use mapreduce.job.user.name

2017-11-03 13:33:22,455 INFO [main] Configuration.deprecation: mapred.reduce.tasks is deprecated. Instead, use mapreduce.job.reduces

2017-11-03 13:33:22,455 INFO [main] Configuration.deprecation: mapred.cache.files.filesizes is deprecated. Instead, use mapreduce.job.cache.files.filesizes

2017-11-03 13:33:22,455 INFO [main] Configuration.deprecation: mapred.output.key.class is deprecated. Instead, use mapreduce.job.output.key.class

2017-11-03 13:33:23,096 INFO [main] mapreduce.JobSubmitter: Submitting

```

tokens for job: job_1509695509797_0001
2017-11-03 13:33:24,497 INFO [main] impl.YarnClientImpl: Submitted application
application_1509695509797_0001 to ResourceManager at /0.0.0.0:8032
2017-11-03 13:33:24,655 INFO [main] mapreduce.Job: The url to track the job:
http://http://localhost:8088/proxy/application_1509695509797_0001/
2017-11-03 13:33:24,662 INFO [main] mapreduce.Job: Running job:
job_1509695509797_0001
2017-11-03 13:33:45,511 INFO [main] mapreduce.Job: Job
job_1509695509797_0001 running in uber mode : false
2017-11-03 13:33:45,532 INFO [main] mapreduce.Job: map 0% reduce 0%
2017-11-03 13:33:56,885 INFO [main] mapreduce.Job: map 100% reduce 0%
2017-11-03 13:33:57,938 INFO [main] mapreduce.Job: Job
job_1509695509797_0001 completed successfully
2017-11-03 13:33:58,296 ERROR [main] tool.ImportTool: Imported Failed: No
enum constant org.apache.hadoop.mapreduce.JobCounter.MB_MILLIS_MAPS
[acadgild@localhost ~]$

```

```

hadoop fs -ls /sqoopout1
17/11/03 13:37:35 WARN util.NativeCodeLoader: Unable to load native-hadoop
library for your platform... using builtin-java classes where applicable
Found 2 items
-rw-r--r--  1 acadgild supergroup      0 2017-11-03 13:33 /sqoopout1/_SUCCESS
-rw-r--r--  1 acadgild supergroup    51 2017-11-03 13:33
/sqoopout1/part-m-000000
[acadgild@localhost ~]$ hadoop fs -cat /sqoopout1/part-m-000000
17/11/03 13:38:10 WARN util.NativeCodeLoader: Unable to load native-hadoop
library for your platform... using builtin-java classes where applicable
1,acadgild,banglure
2,spring,banglure
3,ibm,cochin

```

SQOOP EXPORT

```

mysql> truncate table company;
Query OK, 0 rows affected (0.00 sec)

```

```

mysql> select * from company;
Empty set (0.00 sec)

```

```

[acadgild@localhost ~]$ sqoop export -m 1 --connect jdbc:mysql://localhost/db1
--username 'root' -P --table 'company' --export-dir '/sqoopout1'
--input-fields-terminated-by ','
Warning: /usr/local/sqoop/./hcatalog does not exist! HCatalog jobs will fail.

```

Please set \$HCAT_HOME to the root of your HCatalog installation.
Warning: /usr/local/sqoop/./accumulo does not exist! Accumulo imports will fail.
Please set \$ACCUMULO_HOME to the root of your Accumulo installation.
Warning: /usr/local/sqoop/./zookeeper does not exist! Accumulo imports will fail.
Please set \$ZOOKEEPER_HOME to the root of your Zookeeper installation.
2017-11-03 13:54:36,246 INFO [main] sqoop.Sqoop: Running Sqoop version: 1.4.5
Enter password:
2017-11-03 13:54:42,286 INFO [main] manager.MySQLManager: Preparing to use a MySQL streaming resultset.
2017-11-03 13:54:42,286 INFO [main] tool.CodeGenTool: Beginning code generation
2017-11-03 13:54:42,882 INFO [main] manager.SqlManager: Executing SQL statement: SELECT t.* FROM `company` AS t LIMIT 1
2017-11-03 13:54:42,987 INFO [main] manager.SqlManager: Executing SQL statement: SELECT t.* FROM `company` AS t LIMIT 1
2017-11-03 13:54:43,004 INFO [main] orm.CompilationManager: HADOOP_MAPRED_HOME is /usr/local/hadoop-2.6.0
Note:
/tmp/sqoop-acadgild/compile/a26a4329f42c53071c449566a3c61e25/company.java uses or overrides a deprecated API.
Note: Recompile with -Xlint:deprecation for details.
2017-11-03 13:54:46,815 INFO [main] orm.CompilationManager: Writing jar file: /tmp/sqoop-acadgild/compile/a26a4329f42c53071c449566a3c61e25/company.jar
2017-11-03 13:54:46,860 INFO [main] mapreduce.ExportJobBase: Beginning export of company
SLF4J: Class path contains multiple SLF4J bindings.
SLF4J: Found binding in [jar:file:/usr/local/hbase/lib/slf4j-log4j12-1.6.4.jar!/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: Found binding in [jar:file:/usr/local/hadoop-2.6.0/share/hadoop/common/lib/slf4j-log4j12-1.7.5.jar!/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: See http://www.slf4j.org/codes.html#multiple_bindings for an explanation.
2017-11-03 13:54:47,532 WARN [main] util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
2017-11-03 13:54:47,549 INFO [main] Configuration.deprecation: mapred.jar is deprecated. Instead, use mapreduce.job.jar
2017-11-03 13:54:49,047 INFO [main] Configuration.deprecation: mapred.reduce.tasks.speculative.execution is deprecated. Instead, use mapreduce.reduce.speculative
2017-11-03 13:54:49,083 INFO [main] Configuration.deprecation: mapred.map.tasks.speculative.execution is deprecated. Instead, use mapreduce.map.speculative
2017-11-03 13:54:49,092 INFO [main] Configuration.deprecation: mapred.map.tasks is deprecated. Instead, use mapreduce.job.maps
2017-11-03 13:54:49,509 INFO [main] client.RMProxy: Connecting to

ResourceManager at /0.0.0.0:8032

2017-11-03 13:54:52,027 INFO [main] input.FileInputFormat: Total input paths to process : 1

2017-11-03 13:54:52,053 INFO [main] input.FileInputFormat: Total input paths to process : 1

2017-11-03 13:54:52,212 INFO [main] mapreduce.JobSubmitter: number of splits:1

2017-11-03 13:54:52,246 INFO [main] Configuration.deprecation: mapred.job.name is deprecated. Instead, use mapreduce.job.name

2017-11-03 13:54:52,252 INFO [main] Configuration.deprecation: mapred.cache.files.timestamps is deprecated. Instead, use mapreduce.job.cache.files.timestamps

2017-11-03 13:54:52,252 INFO [main] Configuration.deprecation: mapreduce.map.class is deprecated. Instead, use mapreduce.job.map.class

2017-11-03 13:54:52,252 INFO [main] Configuration.deprecation: mapred.input.dir is deprecated. Instead, use mapreduce.input.fileinputformat.inputdir

2017-11-03 13:54:52,253 INFO [main] Configuration.deprecation: mapred.map.tasks.speculative.execution is deprecated. Instead, use mapreduce.map.speculative

2017-11-03 13:54:52,254 INFO [main] Configuration.deprecation: mapreduce.inputformat.class is deprecated. Instead, use mapreduce.job.inputformat.class

2017-11-03 13:54:52,255 INFO [main] Configuration.deprecation: mapreduce.outputformat.class is deprecated. Instead, use mapreduce.job.outputformat.class

2017-11-03 13:54:52,256 INFO [main] Configuration.deprecation: mapred.cache.files is deprecated. Instead, use mapreduce.job.cache.files

2017-11-03 13:54:52,256 INFO [main] Configuration.deprecation: mapred.working.dir is deprecated. Instead, use mapreduce.job.working.dir

2017-11-03 13:54:52,257 INFO [main] Configuration.deprecation: mapred.mapoutput.value.class is deprecated. Instead, use mapreduce.map.output.value.class

2017-11-03 13:54:52,257 INFO [main] Configuration.deprecation: mapred.mapoutput.key.class is deprecated. Instead, use mapreduce.map.output.key.class

2017-11-03 13:54:52,257 INFO [main] Configuration.deprecation: mapred.job.classpath.files is deprecated. Instead, use mapreduce.job.classpath.files

2017-11-03 13:54:52,257 INFO [main] Configuration.deprecation: user.name is deprecated. Instead, use mapreduce.job.user.name

2017-11-03 13:54:52,257 INFO [main] Configuration.deprecation: mapred.reduce.tasks is deprecated. Instead, use mapreduce.job.reduces

2017-11-03 13:54:52,257 INFO [main] Configuration.deprecation: mapred.cache.files.filesizes is deprecated. Instead, use mapreduce.job.cache.files.filesizes

2017-11-03 13:54:52,592 INFO [main] mapreduce.JobSubmitter: Submitting tokens for job: job_1509695509797_0003

```

2017-11-03 13:54:53,350 INFO [main] impl.YarnClientImpl: Submitted application
application_1509695509797_0003 to ResourceManager at /0.0.0.0:8032
2017-11-03 13:54:53,469 INFO [main] mapreduce.Job: The url to track the job:
http://http://localhost:8088/proxy/application_1509695509797_0003/
2017-11-03 13:54:53,470 INFO [main] mapreduce.Job: Running job:
job_1509695509797_0003
2017-11-03 13:55:08,088 INFO [main] mapreduce.Job: Job
job_1509695509797_0003 running in uber mode : false
2017-11-03 13:55:08,112 INFO [main] mapreduce.Job: map 0% reduce 0%
2017-11-03 13:55:18,971 INFO [main] mapreduce.Job: map 100% reduce 0%
2017-11-03 13:55:20,046 INFO [main] mapreduce.Job: Job
job_1509695509797_0003 completed successfully

```

```
mysql> select * from company;
```

```

+-----+-----+-----+
| id  | name   | location |
+-----+-----+-----+
|  1  | acadgild | banglure |
|  2  | spring  | banglure |
|  3  | ibm     | cochin   |
+-----+-----+-----+
3 rows in set (0.00 sec)

```

How it Works

Sqoop calls the JDBC driver written in the `--connect` statement from the location where Sqoop is installed. The `--username` and `--password` options are used to authenticate the user and Sqoop, internally generates the same command against the MySQL instance.

The `--table` argument defines the MySQL table name, that will receive the data from HDFS. This table must be created prior to running the export command. Sqoop uses the number of columns, their types, and the metadata of the table to validate the data inserted from the HDFS directory. When the export statement is executed, it initiates and creates INSERT statements in MySQL. For example, the `export-job` will read each line of the `input.txt` file from HDFS and produces the following intermediate statements.

2)Transfer data between mysql and hive

```
mysql> grant all on *.* to 'root@localhost' with grant option;
Query OK, 0 rows affected (0.00 sec)
```

```
mysql> flush privileges;
```

Query OK, 0 rows affected (0.00 sec)

mysql> commit;

Query OK, 0 rows affected (0.00 sec)

mysql>

mysql> select * from employee;

id	name	salary
1	vishnu	100
2	vipin	200
3	vijay	300

3 rows in set (0.00 sec)

A Sqoop command is used to transfer table from MySQL to Hive.

```
[acadgild@localhost ~]$ sqoop import --connect jdbc:mysql://localhost/db1  
--username root --table employee --target-dir /myhive --hive-import  
--create-hive-table --hive-table default.Company2Hive -m 1
```

Warning: /usr/local/sqoop/./hcatalog does not exist! HCatalog jobs will fail.
Please set \$HCAT_HOME to the root of your HCatalog installation.

Warning: /usr/local/sqoop/./accumulo does not exist! Accumulo imports will fail.
Please set \$ACCUMULO_HOME to the root of your Accumulo installation.

Warning: /usr/local/sqoop/./zookeeper does not exist! Accumulo imports will fail.
Please set \$ZOOKEEPER_HOME to the root of your Zookeeper installation.

2017-11-03 15:35:53,470 INFO [main] sqoop.Sqoop: Running Sqoop version:
1.4.5

2017-11-03 15:35:53,544 INFO [main] tool.BaseSqoopTool: Using Hive-specific
delimiters for output. You can override

2017-11-03 15:35:53,545 INFO [main] tool.BaseSqoopTool: delimiters with
--fields-terminated-by, etc.

2017-11-03 15:35:54,223 INFO [main] manager.MySQLManager: Preparing to use
a MySQL streaming resultset.

2017-11-03 15:35:54,233 INFO [main] tool.CodeGenTool: Beginning code
generation

2017-11-03 15:35:55,023 INFO [main] manager.SqlManager: Executing SQL
statement: SELECT t.* FROM `employee` AS t LIMIT 1

2017-11-03 15:35:55,126 INFO [main] manager.SqlManager: Executing SQL
statement: SELECT t.* FROM `employee` AS t LIMIT 1

2017-11-03 15:35:55,148 INFO [main] orm.CompilationManager:

HADOOP_MAPRED_HOME is /usr/local/hadoop-2.6.0

Note:

/tmp/sqoop-acadgild/compile/51f32aaf065965adde75af6e8731f58e/employee.jav
a uses or overrides a deprecated API.

Note: Recompile with -Xlint:deprecation for details.

2017-11-03 15:35:59,546 INFO [main] orm.CompilationManager: Writing jar file:

/tmp/sqoop-acadgild/compile/51f32aaf065965adde75af6e8731f58e/employee.jar
2017-11-03 15:35:59,596 WARN [main] manager.MySQLManager: It looks like you
are importing from mysql.
2017-11-03 15:35:59,596 WARN [main] manager.MySQLManager: This transfer
can be faster! Use the --direct
2017-11-03 15:35:59,596 WARN [main] manager.MySQLManager: option to
exercise a MySQL-specific fast path.
2017-11-03 15:35:59,596 INFO [main] manager.MySQLManager: Setting zero
DATETIME behavior to convertToNull (mysql)
2017-11-03 15:35:59,637 INFO [main] mapreduce.ImportJobBase: Beginning
import of employee
SLF4J: Class path contains multiple SLF4J bindings.
SLF4J: Found binding in
[jar:file:/usr/local/hbase/lib/slf4j-log4j12-1.6.4.jar!/org/slf4j/impl/StaticLoggerBinder
.class]
SLF4J: Found binding in
[jar:file:/usr/local/hadoop-2.6.0/share/hadoop/common/lib/slf4j-log4j12-1.7.5.jar!/o
rg/slf4j/impl/StaticLoggerBinder.class]
SLF4J: See http://www.slf4j.org/codes.html#multiple_bindings for an explanation.
2017-11-03 15:36:00,343 WARN [main] util.NativeCodeLoader: Unable to load
native-hadoop library for your platform... using builtin-java classes where
applicable
2017-11-03 15:36:00,396 INFO [main] Configuration.deprecation: mapred.jar is
deprecated. Instead, use mapreduce.job.jar
2017-11-03 15:36:01,744 INFO [main] Configuration.deprecation:
mapred.map.tasks is deprecated. Instead, use mapreduce.job.maps
2017-11-03 15:36:02,141 INFO [main] client.RMProxy: Connecting to
ResourceManager at /0.0.0.0:8032
2017-11-03 15:36:05,755 INFO [main] db.DBInputFormat: Using read committed
transaction isolation
2017-11-03 15:36:05,989 INFO [main] mapreduce.JobSubmitter: number of
splits:1
2017-11-03 15:36:06,019 INFO [main] Configuration.deprecation:
mapred.job.name is deprecated. Instead, use mapreduce.job.name
2017-11-03 15:36:06,020 INFO [main] Configuration.deprecation:
mapred.cache.files.timestamps is deprecated. Instead, use
mapreduce.job.cache.files.timestamps
2017-11-03 15:36:06,020 INFO [main] Configuration.deprecation:
mapreduce.map.class is deprecated. Instead, use mapreduce.job.map.class
2017-11-03 15:36:06,025 INFO [main] Configuration.deprecation:
mapreduce.inputformat.class is deprecated. Instead, use
mapreduce.job.inputformat.class
2017-11-03 15:36:06,025 INFO [main] Configuration.deprecation:
mapreduce.outputformat.class is deprecated. Instead, use
mapreduce.job.outputformat.class
2017-11-03 15:36:06,025 INFO [main] Configuration.deprecation:
mapred.output.value.class is deprecated. Instead, use
mapreduce.job.output.value.class

```

2017-11-03 15:36:06,025 INFO [main] Configuration.deprecation:
mapred.output.dir is deprecated. Instead, use
mapreduce.output.fileoutputformat.outputdir
2017-11-03 15:36:06,025 INFO [main] Configuration.deprecation:
mapred.cache.files is deprecated. Instead, use mapreduce.job.cache.files
2017-11-03 15:36:06,025 INFO [main] Configuration.deprecation:
mapred.working.dir is deprecated. Instead, use mapreduce.job.working.dir
2017-11-03 15:36:06,026 INFO [main] Configuration.deprecation:
mapred.job.classpath.files is deprecated. Instead, use
mapreduce.job.classpath.files
2017-11-03 15:36:06,026 INFO [main] Configuration.deprecation: user.name is
deprecated. Instead, use mapreduce.job.user.name
2017-11-03 15:36:06,026 INFO [main] Configuration.deprecation:
mapred.reduce.tasks is deprecated. Instead, use mapreduce.job.reduces
2017-11-03 15:36:06,026 INFO [main] Configuration.deprecation:
mapred.cache.files.filesizes is deprecated. Instead, use
mapreduce.job.cache.files.filesizes
2017-11-03 15:36:06,026 INFO [main] Configuration.deprecation:
mapred.output.key.class is deprecated. Instead, use
mapreduce.job.output.key.class
2017-11-03 15:36:06,487 INFO [main] mapreduce.JobSubmitter: Submitting
tokens for job: job_1509695509797_0004
2017-11-03 15:36:07,127 INFO [main] impl.YarnClientImpl: Submitted application
application_1509695509797_0004 to ResourceManager at /0.0.0.0:8032
2017-11-03 15:36:07,198 INFO [main] mapreduce.Job: The url to track the job:
http://http://localhost:8088/proxy/application_1509695509797_0004/
2017-11-03 15:36:07,200 INFO [main] mapreduce.Job: Running job:
job_1509695509797_0004
2017-11-03 15:36:23,750 INFO [main] mapreduce.Job: Job
job_1509695509797_0004 running in uber mode : false
2017-11-03 15:36:23,753 INFO [main] mapreduce.Job: map 0% reduce 0%
2017-11-03 15:36:36,412 INFO [main] mapreduce.Job: map 100% reduce 0%
2017-11-03 15:36:37,490 INFO [main] mapreduce.Job: Job
job_1509695509797_0004 completed successfully

```

```
hive> select * from employee;
```

```
OK
```

```

1      vishnu  100
2      vipin  200
3      vijay   300

```

```
Time taken: 3.942 seconds, Fetched: 3 row(s)
```

A Sqoop command is used to transfer selected columns from MySQL to Hive.

Importing columns name and id

```
[acadgild@localhost ~]$ sqoop import --connect jdbc:mysql://localhost/db1
--username 'root' -P --split-by id --columns id,name, --table 'employee'
--target-dir /myhive5 --hive-import --create-hive-table --hive-table
default.Company1Hive -m 1
```

Warning: /usr/local/sqoop/./hcatalog does not exist! HCatalog jobs will fail.
Please set \$HCAT_HOME to the root of your HCatalog installation.

Warning: /usr/local/sqoop/./accumulo does not exist! Accumulo imports will fail.
Please set \$ACCUMULO_HOME to the root of your Accumulo installation.

Warning: /usr/local/sqoop/./zookeeper does not exist! Accumulo imports will fail.
Please set \$ZOOKEEPER_HOME to the root of your Zookeeper installation.

2017-11-03 16:14:26,466 INFO [main] sqoop.Sqoop: Running Sqoop version:
1.4.5

Enter password:

2017-11-03 16:14:27,903 INFO [main] tool.BaseSqoopTool: Using Hive-specific
delimiters for output. You can override

2017-11-03 16:14:27,903 INFO [main] tool.BaseSqoopTool: delimiters with
--fields-terminated-by, etc.

2017-11-03 16:14:28,383 INFO [main] manager.MySQLManager: Preparing to use
a MySQL streaming resultset.

2017-11-03 16:14:28,409 INFO [main] tool.CodeGenTool: Beginning code
generation

2017-11-03 16:14:29,105 INFO [main] manager.SqlManager: Executing SQL
statement: SELECT t.* FROM `employee` AS t LIMIT 1

2017-11-03 16:14:29,260 INFO [main] manager.SqlManager: Executing SQL
statement: SELECT t.* FROM `employee` AS t LIMIT 1

2017-11-03 16:14:29,280 INFO [main] orm.CompilationManager:
HADOOP_MAPRED_HOME is /usr/local/hadoop-2.6.0

Note:

/tmp/sqoop-acadgild/compile/37a2c328df810779e9a15c92924129a3/employee.ja
va uses or overrides a deprecated API.

Note: Recompile with -Xlint:deprecation for details.

2017-11-03 16:14:33,422 INFO [main] orm.CompilationManager: Writing jar file:
/tmp/sqoop-acadgild/compile/37a2c328df810779e9a15c92924129a3/employee.ja
r

2017-11-03 16:14:33,499 WARN [main] manager.MySQLManager: It looks like you
are importing from mysql.

2017-11-03 16:14:33,499 WARN [main] manager.MySQLManager: This transfer
can be faster! Use the --direct

2017-11-03 16:14:33,499 WARN [main] manager.MySQLManager: option to
exercise a MySQL-specific fast path.

2017-11-03 16:14:33,499 INFO [main] manager.MySQLManager: Setting zero
DATETIME behavior to convertToNull (mysql)

2017-11-03 16:14:33,503 INFO [main] mapreduce.ImportJobBase: Beginning
import of employee

SLF4J: Class path contains multiple SLF4J bindings.

SLF4J: Found binding in

[jar:file:/usr/local/hbase/lib/slf4j-log4j12-1.6.4.jar!/org/slf4j/impl/StaticLoggerBinder

.class]
SLF4J: Found binding in
[jar:file:/usr/local/hadoop-2.6.0/share/hadoop/common/lib/slf4j-log4j12-1.7.5.jar!/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: See http://www.slf4j.org/codes.html#multiple_bindings for an explanation.
2017-11-03 16:14:34,256 WARN [main] util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
2017-11-03 16:14:34,273 INFO [main] Configuration.deprecation: mapred.jar is deprecated. Instead, use mapreduce.job.jar
2017-11-03 16:14:35,632 INFO [main] Configuration.deprecation: mapred.map.tasks is deprecated. Instead, use mapreduce.job.maps
2017-11-03 16:14:36,031 INFO [main] client.RMProxy: Connecting to ResourceManager at /0.0.0.0:8032
2017-11-03 16:14:39,134 INFO [main] db.DBInputFormat: Using read committed transaction isolation
2017-11-03 16:14:39,421 INFO [main] mapreduce.JobSubmitter: number of splits:1
2017-11-03 16:14:39,502 INFO [main] Configuration.deprecation: mapred.job.name is deprecated. Instead, use mapreduce.job.name
2017-11-03 16:14:39,505 INFO [main] Configuration.deprecation: mapred.cache.files.timestamps is deprecated. Instead, use mapreduce.job.cache.files.timestamps
2017-11-03 16:14:39,505 INFO [main] Configuration.deprecation: mapreduce.map.class is deprecated. Instead, use mapreduce.job.map.class
2017-11-03 16:14:39,507 INFO [main] Configuration.deprecation: mapreduce.inputformat.class is deprecated. Instead, use mapreduce.job.inputformat.class
2017-11-03 16:14:39,513 INFO [main] Configuration.deprecation: mapreduce.outputformat.class is deprecated. Instead, use mapreduce.job.outputformat.class
2017-11-03 16:14:39,519 INFO [main] Configuration.deprecation: mapred.output.value.class is deprecated. Instead, use mapreduce.job.output.value.class
2017-11-03 16:14:39,519 INFO [main] Configuration.deprecation: mapred.output.dir is deprecated. Instead, use mapreduce.output.fileoutputformat.outputdir
2017-11-03 16:14:39,519 INFO [main] Configuration.deprecation: mapred.cache.files is deprecated. Instead, use mapreduce.job.cache.files
2017-11-03 16:14:39,520 INFO [main] Configuration.deprecation: mapred.working.dir is deprecated. Instead, use mapreduce.job.working.dir
2017-11-03 16:14:39,520 INFO [main] Configuration.deprecation: mapred.job.classpath.files is deprecated. Instead, use mapreduce.job.classpath.files
2017-11-03 16:14:39,520 INFO [main] Configuration.deprecation: user.name is deprecated. Instead, use mapreduce.job.user.name
2017-11-03 16:14:39,524 INFO [main] Configuration.deprecation: mapred.reduce.tasks is deprecated. Instead, use mapreduce.job.reduces

2017-11-03 16:14:39,524 INFO [main] Configuration.deprecation:
mapred.cache.files.filesizes is deprecated. Instead, use
mapreduce.job.cache.files.filesizes
2017-11-03 16:14:39,524 INFO [main] Configuration.deprecation:
mapred.output.key.class is deprecated. Instead, use
mapreduce.job.output.key.class
2017-11-03 16:14:39,964 INFO [main] mapreduce.JobSubmitter: Submitting
tokens for job: job_1509695509797_0007
2017-11-03 16:14:40,797 INFO [main] impl.YarnClientImpl: Submitted application
application_1509695509797_0007 to ResourceManager at /0.0.0.0:8032
2017-11-03 16:14:40,948 INFO [main] mapreduce.Job: The url to track the job:
http://localhost:8088/proxy/application_1509695509797_0007/
2017-11-03 16:14:40,953 INFO [main] mapreduce.Job: Running job:
job_1509695509797_0007
2017-11-03 16:14:57,587 INFO [main] mapreduce.Job: Job
job_1509695509797_0007 running in uber mode : false
2017-11-03 16:14:57,595 INFO [main] mapreduce.Job: map 0% reduce 0%
2017-11-03 16:15:09,046 INFO [main] mapreduce.Job: map 100% reduce 0%
2017-11-03 16:15:10,144 INFO [main] mapreduce.Job: Job
job_1509695509797_0007 completed successfully
2017-11-03 16:15:10,525 ERROR [main] tool.ImportTool: Imported Failed: No
enum constant org.apache.hadoop.mapreduce.JobCounter.MB_MILLIS_MAPS
[acadgild@localhost ~]\$